Burry Is Right About the GPUs. He's Wrong About NVIDIA.

The GPU "values" he attacked are really rent. But the shortage propping up rent outlasts his puts, and if rent does fall, it hits the neoclouds before it reaches NVIDIA.



PRIMARY SOURCE ARBITRAGE

OCT 05, 2026

The short version

- **What happened:** On October 1, Michael Burry argued that NVIDIA's chart of GPU resale values is a rent forecast dressed up as a price. In late September, he had replaced his NVIDIA short with September 2027 puts struck in the mid-$100s.
- **Where he's right:** Silicon Data, the firm behind the chart, describes the values in its own methodology as discounted future rent. They rise and fall with rent, and the last time a GPU shortage eased, H100 rent fell 44% in five months.
- **Where I disagree:** The memory shortage holding rents up is backed by the people paying for it, not just the people selling it, and it runs through 2028. On my arithmetic, B200 rent would have to fall about 65% before the chip fails to pay for itself inside six years. And the lessors Burry is worried about, mostly pay off their loans inside their customer contracts, so they don't need to re-rent old chips to survive.
- **What I hold:** NVIDIA.

- **What would prove me wrong:** Silicon Data's B200 rent index below $2.85 for three straight monthly readings before September 2027, roughly where a new B200 stops paying for itself in six years at 50% utilization, or at 70% with money costing 10% a year. I'd start reviewing below $3.93, where half-empty clusters funded at 10% stop paying. Or any of Micron, Samsung or SK hynix saying memory supply catches up with demand in calendar 2027.

The Big Short

The first movie I watched when I started investing was *The Big Short*. What stayed with me was Michael Burry alone in his office reading the mortgage-bond prospectuses nobody else had bothered to open, then trusting what he found over what everyone around him was saying. I built this publication around that habit.

Since then I've watched him make bear call after bear call on the AI trade, and I mostly scrolled past. On October 1, he published "Don't Believe Your Lyin' Eyes," aimed squarely at NVIDIA. In late September, according to Stocktwits, he had covered his short positions in Micron, Nebius, CoreWeave, Caterpillar, Palantir, the iShares Semiconductor ETF and NVIDIA, and replaced the NVIDIA short with September 2027 put options struck in the mid-$100s.

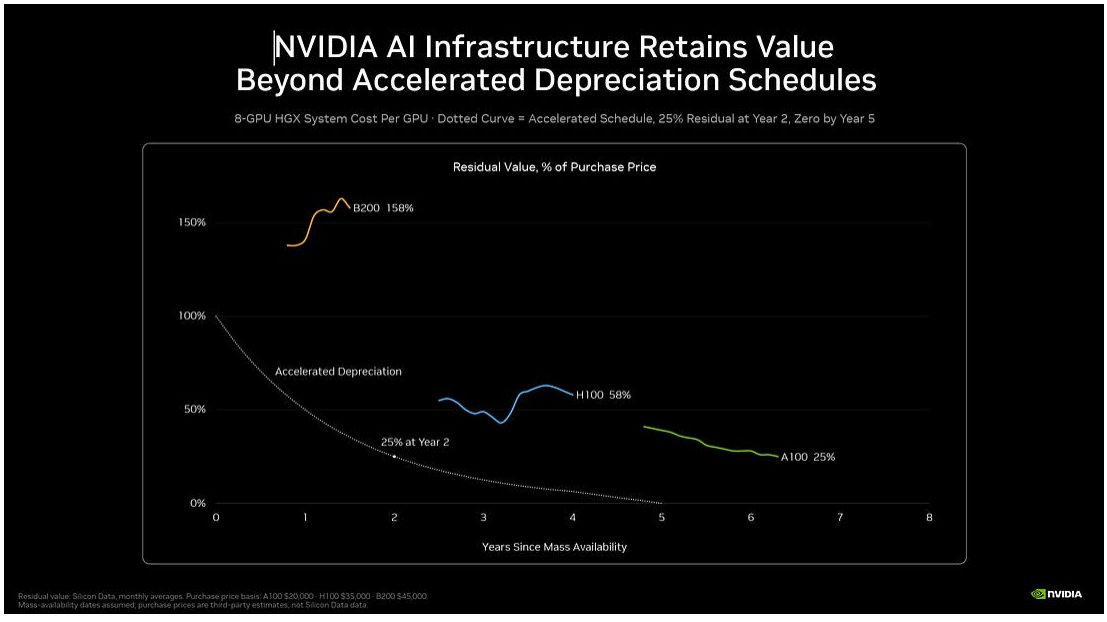
NVIDIA closed up 1.1% that day at $230.86. The next day it traded as high as $237.88, a new 52-week high, before closing at $233.95. For puts struck around $150 to pay off at expiry, the stock has to fall at least 36% from there.

I own NVIDIA, so I'm on the other side of his trade. After months of scrolling past him, I thought he deserved a proper reading.

What he’s actually arguing

The target is a single slide. In September, NVIDIA took a deck on the road to investors, and page 22 is titled “NVIDIA AI Infrastructure Retains Value Beyond Accelerated Depreciation Schedules.” It plots what three generations of NVIDIA GPUs are supposedly still worth, as a percentage of their purchase price, against a dotted line labelled “Accelerated Schedule, 25% Residual at Year 2, Zero by Year 5.”

The three curves end far above the dotted line. The B200 is shown at 158% of its purchase price after about a year and a half, the H100 at 58% after four years, and the A100 at 25% after six. The footnote reads: “Residual value: Silicon Data, monthly averages. Purchase price basis: A100 $20,000 · H100 $35,000 · B200 $45,000. Mass-availability dates assumed; purchase prices are third-party estimates, not Silicon Data data.”



NVIDIA's residual-value slide: A100, H100 and B200 values against a five-year accelerated depreciation schedule. Source: NVIDIA Non-Deal Roadshow presentation, September 2026, p.22.

The message is that the buyers of these chips are writing them off too fast. A lot of NVIDIA bulls took it as the end of the depreciation debate.

Burry's post is paywalled, and I haven't read the paid half. From the free section and the press summaries, his argument has three parts.

1. The values on the chart aren't what anyone paid for a used GPU. They come from a discounted cash flow of projected rent over an assumed eight-year life, so comparing them against a five-year write-off curve is apples to oranges.
2. Older GPUs are renting well because memory and power are scarce, not because the chips are durable.
3. And the financing growing up around these chips looks like the computer-leasing boom of the late 1960s.

The free half is mostly that last part, told through a scene from *The Money Game*, published in 1968 under the pen name Adam Smith (the writer George Goodman, not the economist). A speculator called the Great Winfield introduces the author to his young traders, "Billy the Kid, Johnny the Kid, and Sheldon the Kid." Billy has turned $5,000 into half a million in six months buying the convertible bonds of computer-leasing companies. Asked whether he's borrowing, he says, "Not too heavily, sir! I put up at least three percent cash. When I am conservative, I put up five percent cash." His pitch is that "The need for computers is practically infinite... The surface has barely been scratched." When the author asks how fast the computers are being written off, Winfield laughs at him. Questions like that, he says, "show you're middle-aged."

Burry casts himself as the author and NVIDIA's bulls as the Kids. As the history is usually told, the leasing companies spread the cost of IBM machines over far more years than IBM itself did, and their stocks collapsed once IBM's System/370 made the older machines hard to rent out.

His sharpest line, as quoted by Benzinga, is this:

> **“A chip earning more than it costs is a sign of scarcity due to memory and power shortages, not GPU die shortages and not durability.”**

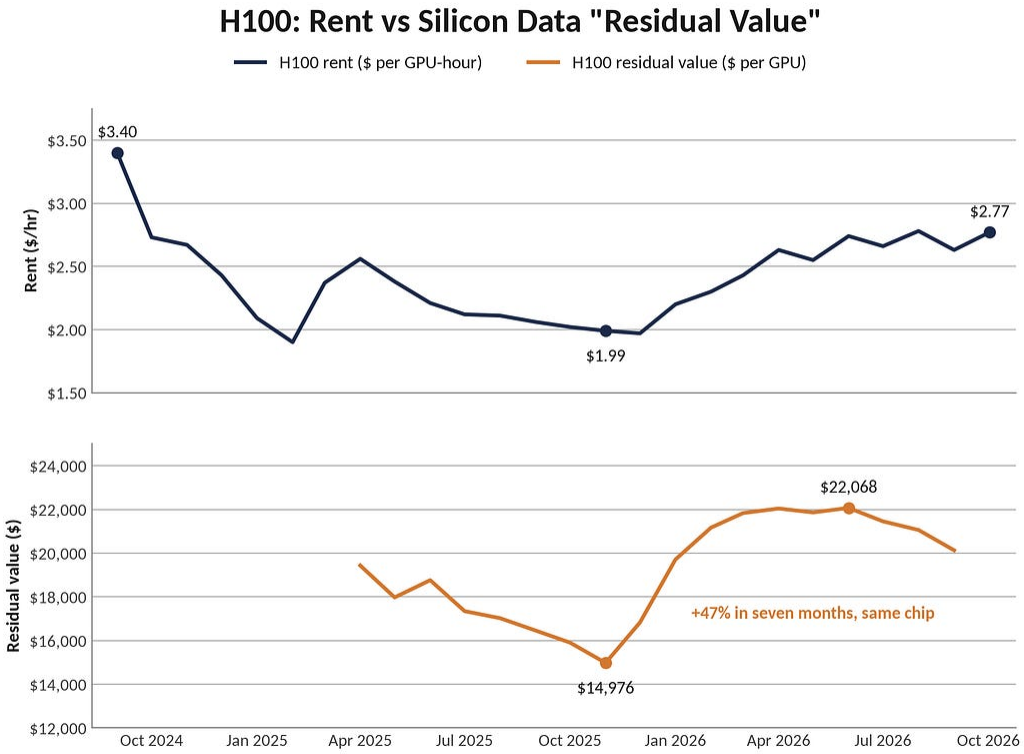
Where he’s right

I started with the footnote, because that’s where the argument lives. Silicon Data publishes a one-page methodology for its residual values, dated June 29, 2026. The values come from a discounted cash flow: an estimate of what something is worth today based on the money it should earn in the future. The model projects net rent over the chip’s remaining life, which it sets at eight years from the date the chip became widely available. Rental rates come from Silicon Data’s own forward curve, its projection of rent for each of the next 36 months, built from today’s prices for rental contracts of one to 36 months. Beyond 36 months, the model holds the last rate flat and applies a decay it doesn’t disclose.

So Burry describes the method accurately. Nobody sold a used B200 for $72,695, which is Silicon Data’s value for one on September 30 and 162% of the $45,000 price on NVIDIA’s slide. That figure is what a B200 should earn over the rest of its eight-year life if rent roughly holds. At the current B200 index of $5.86 an hour, a chip rented 70% of the time brings in about $36,000 a year. Silicon Data’s model adds up almost seven more years of rent like that, after costs and a discount rate it doesn’t publish.

The clearest evidence is how the value moves. Silicon Data’s H100 residual value was $14,976 in November 2025 and $22,068 by June 2026, up 47% in seven months on the same piece of silicon. The chip didn’t get any better. Rent did. Over the same stretch, Silicon Data’s price for a three-year H100 rental contract, which its forward curve is

built from, rose from $1.53 to $2.26 an hour, about 48%. The “value” is today’s rental prices, projected forward.



H100 rental index and H100 residual value, monthly, September 2024 to October 2026. Rent went from $1.99 an hour in November 2025 to $2.74 in June 2026. These are first-of-month readings, so the daily low of about $1.96 in late November doesn't show. Source: Silicon Data (SDH100RT; GPU Residual Value), accessed October 2, 2026.

NVIDIA makes the opposite case with the same numbers. In an August 11 post on NVIDIA’s blog, CEO Jensen Huang cites one-year H100 rental pricing rising “from about $1.70 per GPU-hour in October 2025 to about $2.35 per GPU-hour in March 2026” as evidence of “the durability of NVIDIA compute economics.” Burry reads that rise as scarcity. The data can’t settle which of them is right until the shortage ends.

The people renting these chips out half-concede Burry's scarcity point. At Goldman Sachs' Communacopia conference on September 8, Nebius's Chief Revenue Officer, Marc Boroditsky, said:

> **"So that might lead you to believe that older chips are being used because people have no choice. There's a little bit of that."**

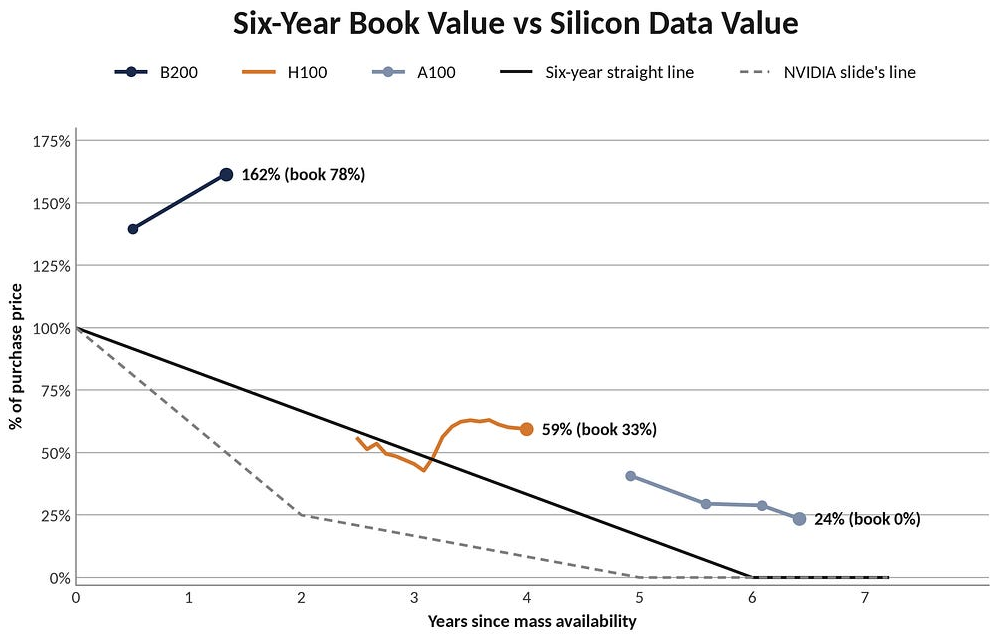
He went on to argue that most older-chip workloads suit those chips, and I think that's fair. Still, "a little bit of that" is a seller telling you scarcity is in the price. CoreWeave's CFO, Nitin Agrawal, on the August 11 earnings call, called installed capacity "a scarce, valuable asset" in "a market where new capacity is supply constrained."

We've also seen what happens when a shortage eases. Silicon Data's H100 index was $3.40 an hour on September 1, 2024. By February 1, 2025, it was $1.90, down 44% in five months, as the first H100 shortage ended and new supply arrived. It bounced to $2.56 in April and then slid back to about $1.96 by late November. **If Burry had made this argument in 2024, he'd have been right.**

There's one more thing about the chart that bothers me, though not for Burry's reason. The dotted line, worth 25% after two years and nothing after five, is harsher than how the buyers actually book these chips. CoreWeave's S-1 shows it moved the useful life of its data-center computing equipment from five years to six, effective January 1, 2023, and writes it off in equal annual amounts. NVIDIA's own deck, one page earlier, uses "$60M/MW build · 6-yr depreciation" for its payback example. Then it switches to the harsher curve for the comparison that flatters it most.

The odd part is that NVIDIA didn't need to. I redrew the comparison against a plain six-year straight line, where a chip's book value falls by a sixth every year. I took each chip's age from Silicon Data's own dates (May 2020 for the A100, October 2022 for the H100, June 2025 for the B200). On September 30, the B200 was worth 162% of its price against

a book value of 78%, the H100 59% against 33%, and the A100 24% against nothing.



Silicon Data value as % of purchase price by chip age, against a six-year straight line and NVIDIA's accelerated schedule. Values as of September 30; NVIDIA's slide uses monthly averages, which is why it shows 158%, 58% and 25%. Source: Silicon Data; NVIDIA Non-Deal Roadshow presentation, September 2026, p.22; author's calculations.

All three sit above the line today. The H100 didn't always: from April to November 2025, while its rent was sliding, its value sat below the six-year line, and it only climbed back above it when the memory shortage pushed rent up. So the argument isn't really about NVIDIA's choice of line. It's about whether those curves hold, and the curves are rent. That's the part of Burry's case I accept. Where I part ways with him is on when he needs rent to fall, and on who he thinks pays when it does.

His clock

Burry's catalyst is the bottleneck easing: once memory and power stop being scarce, a wave of newer chips arrives, and old-chip rent falls. The memory makers say that doesn't happen before 2028. They would say that, since they're the ones selling into the shortage, so I wanted evidence from people who'd rather it weren't true.

The best witness is the biggest buyer. On NVIDIA's earnings call on August 26, CFO Colette Kress said:

> **"As you are already aware, we are experiencing extreme pricing conditions in memory. The magnitude of the price increase has exceeded our prior expectations and are headed even higher into next year."**

She cut NVIDIA's gross margin guidance to a 71% to 72% trough in the fourth quarter because of it, and added that "we expect supply to remain a bottleneck at least through the end of fiscal year '28," which ends in January 2028. **A company doesn't lower its own margin guidance to talk up a supplier.**

Then there's money, which is harder to fake than words. NVIDIA's purchase commitments rose from $119 billion to $279 billion in a single quarter, "primarily related to the procurement of memory," according to its September investor deck. By fiscal year, the commitments run $92 billion for the rest of fiscal 2027, $87 billion for fiscal 2028 and $88 billion for fiscal 2029, then drop to $6 billion in fiscal 2030. Micron's customers have put up $32 billion of financial commitments under its long-term agreements, "the vast majority of which are cash deposits," and SK hynix and Samsung describe the same thing in their own contracts. Customers don't prepay billions for something they expect to get cheaper.

The third witness is the construction calendar, which nobody can talk up. Samsung said on July 30 that the lead time from new fab construction to actual wafer production “exceeds 3 years.” Micron’s dates on its September 30 call: HBM packaging in Singapore in early 2027, its first Idaho fab and Taiwan’s Tongluo site in mid-2027, a second Idaho fab and the Japan expansion in late 2028, New York in 2030. Micron added that production from a new fab “becomes more meaningful starting a few quarters after initial output.” **Anything not already being built today can’t produce meaningful memory before Burry’s puts expire.**

Only then do the sellers’ own forecasts carry weight. Micron’s CEO, Sanjay Mehrotra, on September 30:

> **“We expect memory and storage supply-demand conditions to be much tighter in calendar 2027 and 2028 than they were in 2026.”**

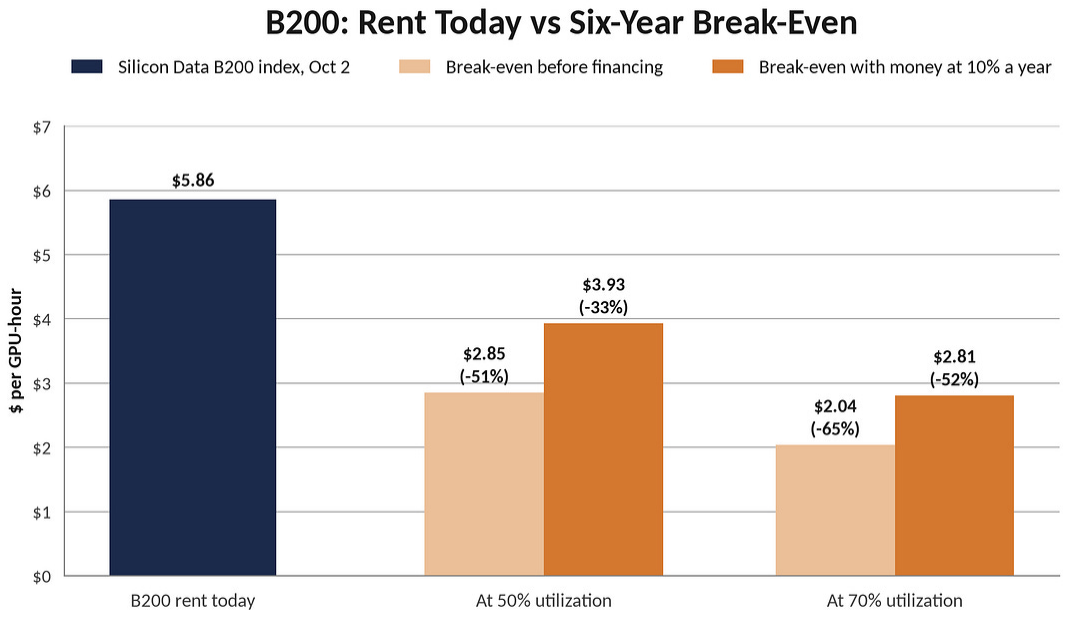
He added that Micron does “not have line of sight to when supply and demand will return to balance.” Samsung, on July 30:

> **“We believe it will be unlikely to see any significant increase in incremental supply through 2028.”**

The dissent is in the stock market, which values Micron at 6.1x NTM P/E and Samsung at 4.2x (LSEG Workspace, October 2), the way markets price a business near its peak. Samsung itself says “Beyond 2029, it is hard to say,” so I read the dissent as a bet on what happens after 2028, not on September 2027. Even the market's skeptics aren't on Burry's clock.

That leaves the question of how far rent can fall before the lessors are actually in trouble. Silicon Data doesn't publish its utilization or cost assumptions, so I used my own: a chip rented 70% of the time, with power, staff, and the building eating 40% of the rent. That's close to CoreWeave's own numbers; its adjusted EBITDA margin, which leaves

out the cost of the chips themselves, was 59% in the second quarter. On the $45,000 per-GPU system cost on NVIDIA's slide (a third-party estimate), today's $5.86 an hour earns the chip back in about 2.1 years.



B200 rent today vs the rent needed to earn back $45,000 in six years, at 50% and 70% utilization, before financing and with money costing 10% a year. Source: Silicon Data B200 index, October 2, 2026; author's calculations.

On my numbers, B200 rent would have to fall about 65% to roughly $2 an hour, before the chip stops paying for itself within the six years its buyers depreciate it over. That assumes 70% utilization, which is conservative today: Kress said in August that “NVIDIA compute is fully utilized across every cloud we serve,” and CoreWeave and Nebius both say they’re sold out. But utilization is the first thing to fall in a glut, and the money behind the chip isn’t free. The chart above shows how fast the cushion shrinks:

- At 50% utilization, a 51% fall is enough ($2.85).
- With money costing 10% a year, about what CoreWeave’s riskiest loan costs, a 52% fall is enough ($2.81).

- With both, half-empty clusters funded at 10%, a fall of a third is enough ($3.93).

That last case is the one that would make me wrong, which is why rule 6 starts a review at $3.93. The worst fall in the H100 index so far was 44%, and it happened while supply was flooding in.

Burry needs a fall that size, sooner, and at a time when the memory makers and their largest customer both say supply stays tight.

To be fair to his trade, the puts don't strictly need rent to fall. NVIDIA trades at about 16.5x NTM P/E (LSEG). If the market starts pricing the end of the shortage early, the way it already prices Micron at 6.1x, the stock could fall 36% with rent intact. That's the real risk to my position, and it's a bet on the multiple rather than the GPU math in his post. Price isn't on my ledger by design. If the stock fell 36% with every rule below intact, I'd hold and add, because the business would be doing what I bought it for and only the price would have changed. Whether the multiple holds comes down to what the market thinks demand looks like after the shortage: whether cheaper tokens keep pulling in more use, and whether newer workloads like robotics arrive in time. I think that's the part of the story analysts' models capture worst, and I come back to it at the end.

Who actually holds the risk?

Burry's 1968 story had two kinds of borrowers. Billy the Kid borrowed to buy the leasing stocks. The leasing companies borrowed to buy IBM machines, and they needed to keep renting those machines out for years to pay the loans back. Today's equivalent of the leasing companies are the neoclouds, which buy NVIDIA GPUs, mostly with borrowed money, and rent out the computing time. So the test is whether their loans depend on re-renting old chips.

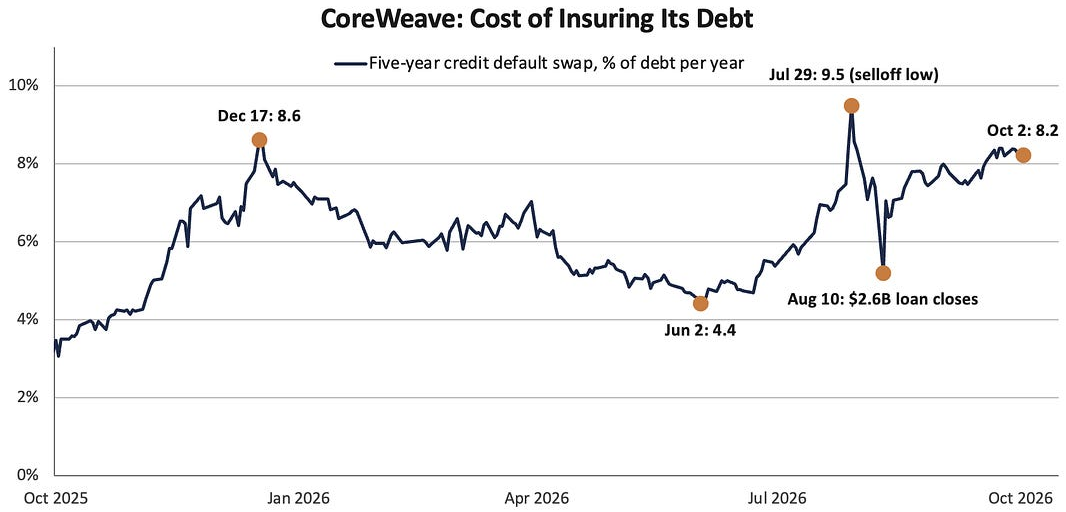
At Nebius, the money mostly comes back inside the contract. On its August 12 earnings call, Nebius said its core deals run one to three years at $20 million to $25 million of revenue per megawatt, with “less than 2-year payback,” and that customers prepay 50% to 60% of the associated capital spending. About 70% of the deals it closed in the second quarter included a prepayment. Its first asset-backed loan, $775 million in July, was secured by deployed GPUs and an investment-grade customer’s contract and priced at SOFR plus 2.50%. SOFR is the overnight rate banks charge each other against Treasuries; it was 3.87% on October 1, so that loan costs Nebius about 6.4%. Nebius says the loan and the cash coming in under that customer’s contract together cover “more than 100% of the capital expenditure” needed to deploy the chips. A one-year deal with a two-year payback still needs renewing, which is why “mostly” is the honest word.

CoreWeave’s main loans sit in per-contract special-purpose companies, which are repaid from that customer’s payments before CoreWeave sees anything. CFO Nitin Agrawal said on August 11 that each is designed so that “when an initial contract ends, the cluster no longer has any leverage.” Re-renting after that is upside, not survival.

There’s one exception, and it’s exactly the shape Burry is describing. On August 10, CoreWeave closed a $2.6 billion loan it calls DDTL 5.5. It runs about five years, while the customer contracts behind it average about three. Think of a five-year mortgage on a house whose tenant signed a three-year lease. For the last two years, the lenders get paid only if CoreWeave finds new tenants for those GPUs at a decent rent. CoreWeave’s own press release, filed with the SEC, says the lenders are “signalling confidence in the long-term value of NVIDIA GPUs running on CoreWeave’s cloud platform and a willingness to underwrite renewal risk.” The lenders charged for it. The loan priced at SOFR plus 5.50%, a full percentage point wider than CoreWeave’s own previous loan of this kind three months earlier, at the same credit ratings, and lenders paid

about $96 for every $100 they'll be repaid. At today's SOFR, that works out to roughly 10.4% a year if the loan runs its full five years.

That price is where 2026 differs from Burry's other famous trade. In 2006, mortgage risk was underpriced: bonds backed by subprime home loans carried investment-grade ratings, and insuring them against default cost very little. Here, the loan carrying renewal risk is rated below investment grade (Ba2 at Moody's, BB+ at Fitch) and priced like it, about three percentage points more than Nebius pays on contracted cash flows. Insuring CoreWeave's debt against default for five years, through its credit default swap, cost about 8.2% a year on October 2 (Markit, via LSEG), about two and a half times what it cost a year earlier. Under the standard pricing model, which assumes lenders get back 40 cents on the dollar, that implies roughly even odds of default over five years (LSEG's model puts it at 52%). That overstates the real odds, since insurance always costs more than the expected loss, but it's the opposite of 2006's complacency. **The credit market isn't hiding the risk.** It's charging for it, mostly at the company level rather than inside the contract-backed loans.



CoreWeave five-year credit default swap spread, daily, September 30, 2025 to October 2, 2026. Source: Markit via LSEG.

The one place the contract-backed structure is weaker than it looks is the customers themselves. JPMorgan's analysts made the point in their September 24 upgrade:

> **"Several major counterparties to CoreWeave fund their own AI investment from capital markets rather than operating cash flow, so their ability to honor multi-year commitments is financing-sensitive."**

A loan backed by Microsoft's contract is one thing. A loan that looks through to an AI lab that is itself raising money is another, and I'll come back to it at the end.

Billy the Kid put up three percent cash. Nebius's big customers prepay fifty to sixty percent of the cost before the chips are switched on.

The Kid did turn up this year, just not where Burry is looking. Leopold Aschenbrenner's hedge fund, Situational Awareness, had used as much as 400% leverage on AI infrastructure stocks. Its biggest disclosed holdings, Nebius, Sandisk, Micron and CoreWeave, each fell more than 35% in July. Margin calls forced it to sell its public positions to Citadel, with talks starting July 29. CNBC later described that sale as the bottom of the selloff that had started in June. Micron, the iShares Semiconductor ETF and NVIDIA all closed at their lows of the summer on July 29. **The leverage that blew up sat in the stocks, not the GPU loans**, and CoreWeave closed DDTL 5.5 twelve days later.

Who was IBM

In 1968, the leasing companies were hurt by the company that made their machines obsolete. Today that company is NVIDIA. Its September deck says Vera Rubin delivers "up to 30X better AI factory token

throughput and 45X lower cost per million tokens" than the GB300. Every new generation is what pushes older chips down the rent ladder.

That would make NVIDIA the safe party in Burry's story, except for one difference from IBM: NVIDIA has helped finance the lessors. Its 10-Q shows three kinds of exposure.

Guarantees, $108.5 billion at most. $105 billion backs the land, power, and buildings at SB Energy's 4.25-gigawatt Ohio campus, leased to OpenAI for 20 years. The other $3.5 billion backs similar costs for AI cloud partners. NVIDIA pays only if a tenant defaults.

AI cloud agreements, $36 billion. In the 10-Q's words: "if AI clouds do not successfully sell committed capacity to third-party customers, we have agreed to purchase that capacity." Kress called it "a minimum revenue guarantee," with NVIDIA sharing "in a portion of the neoclouds' revenue earned above that floor." The commitments typically run six years and shrink as the clouds rent the capacity to someone else.

Leases, $20 billion. Fifteen-year data-center leases NVIDIA intends to pass on to other companies.

Added up, the worst case is $164.5 billion, about 37% of the $447 billion of EBITDA the LSEG consensus expects NVIDIA to earn in fiscal 2028, and about 3% of its $5.64 trillion market value. (I've left out $29 billion of cloud services NVIDIA buys for its own research.) Almost none of it is live yet. The Ohio guarantees only start when the first building is finished, which NVIDIA expects in fiscal 2029, so **none of the $105 billion is live before Burry's puts expire**, and they end if OpenAI earns a satisfactory credit rating. The AI cloud commitments due in fiscal 2028 come to about $6 billion.

That isn't life-threatening, but it isn't nothing either. A guarantee only gets called if the party behind it can't pay, and the most likely reason OpenAI couldn't pay its Ohio landlord is that its own funding had dried

up. In that world, OpenAI would also be buying fewer GPUs. **The guarantees and the orders fail together.**

There's a fourth exposure that isn't on that list, and it's the closest thing to Burry's bet on NVIDIA's own books. Alongside the $500 billion financing platforms NVIDIA announced on August 10, Huang wrote in the blog post I quoted earlier that NVIDIA "may provide a residual-value support mechanism for up to 25% of an opportunity." In plain terms, NVIDIA would cover part of the shortfall if the chips are worth less than expected when a lease ends, which matters most to lenders whose loans outlast the rental contracts behind them. That's a direct bet that residual values hold. At the full 25% on all $500 billion, it would come to $125 billion, but none of it is signed yet: the partnerships "remain subject to execution of the final agreements."

Huang argues NVIDIA can offer the support because its compute is "fungible, universally adopted, software-upgradable and redeployable," which is the durability claim Burry is disputing. But the Financial Times reported on September 29 that NVIDIA has been talking to insurers about taking some of that risk off its hands, and shared chip-depreciation and compute-price data with at least one of them. **NVIDIA is measuring the same risk Burry is. It just wants insurers to carry part of it.**

So the precise version of the 1968 analogy is this. If the lessors get into trouble, NVIDIA is hurt last and in a particular order. First, nothing happens: it was paid for the chips already shipped. Then the guarantees and revenue floors start costing it cash. Then, if the trouble lasts, the order book takes the hit, because a neocloud that can't borrow can't buy. That's the business NVIDIA calls ACIE, which includes the neoclouds: $40 billion of its $89 billion of data-center revenue last quarter.

We've had one small test already. In late November 2025, H100 rent hit its low of about $1.96 an hour, and in December CoreWeave's default

insurance jumped to about 8.6 points. Over the next two quarters, NVIDIA's YoY revenue growth went from 73% to 85% to 106%, and ACIE revenue grew 25% from the first quarter to the second. Orders are placed quarters ahead of shipment, so this isn't a clean test, but a few months of stress at the lessors didn't show up in the order book. A longer stretch would, and the ledger below says what I'm watching for.

Burry is right that the residual values are rent, and right that scarcity is in the price. **I think his clock is wrong, and if he's right about rent, the damage shows up first at the lessors' company level, where the credit market is already charging for it, not at the company selling them the chips.** What his post circles without naming is what happens to NVIDIA once the shortage does end: where rent settles, whether the neoclouds can upgrade their fleets instead of re-renting old chips, and who keeps paying for the compute. Kress put the AI labs NVIDIA expects to back with its own balance sheet at roughly a quarter of next year's business, and those labs are still raising money to pay their bills. That's next week's post.

Break conditions

Rules 1 to 5 carry forward from my August 28 post, word for word. Rules 6 to 9 are new.

1. Non-GAAP gross margin below 72% for two consecutive quarters during the Vera Rubin ramp. Set Jun 2026. Action: exit review. **Not triggered.** Q2 75.0%; Q3 guided 74%; Q4 guided 71 to 72%. One quarter below 72% is guided, so Q1 FY28 is the test.

2. ACIE/Hyperscale split reversing. Threshold: Hyperscale reaccelerates while ACIE decelerates. Set Jun 2026. Action: exit review. **Not triggered.** ACIE +25% q/q vs Hyperscale +13%.

3. Any hyperscaler explicitly framing custom silicon as primary AI training compute on two consecutive earnings calls combined with a sequential decline in Hyperscale revenue. Set Jun 2026. Action: exit review. **Not triggered.**

4. Vera CPU shipments delayed beyond Q1 FY2028. Set Jun 2026. Action: exit review. **Not triggered.** Shipping to lead partners.

5. FY2028 revenue guide walked back below 60% growth on any subsequent earnings call. Set Aug 28, 2026. Action: exit review. **Not triggered.** 70%.

6. GPU rent. Silicon Data B200 rental index for three consecutive monthly readings before September 2027. Below $3.93, the break-even for a $45,000 B200 at 50% utilization with money costing 10% a year: review, starting with whether the neoclouds are still sold out. Below $2.85, roughly the break-even at 50% utilization before financing, or at 70% with money costing 10%: exit review. Set Oct 5, 2026. Action: review at $3.93, exit review at $2.85. **Not triggered.** $5.86 on Oct 2; its lowest reading in the past year was $4.28 in January, 9% above the review level.

7. Memory supply. Micron, Samsung or SK hynix says memory supply and demand balance in calendar 2027. Set Oct 5, 2026. Action: exit review. **Not triggered.** All three say tighter in 2027.

8. Lessor stress reaches the order book. ACIE revenue declines quarter over quarter. Set Oct 5, 2026. Action: exit review. **Not triggered.** +25% q/q in Q2.

9. Guarantees. NVIDIA records any payment or loss under its AI cloud, land/power/shell, SB Energy or residual-value support arrangements. Set Oct 5, 2026. Action: review and resize. **Not triggered.**

What would change your mind on NVIDIA faster: a fall in GPU rent, or a fall in what the AI labs are willing to spend? I'd like to know which one readers are watching.

Disclosure: I hold a position in NVIDIA. This material is for informational purposes only and does not constitute investment advice. Sources:

- *Michael Burry, "Don't Believe Your Lyin' Eyes, GPU Depreciation & Useful Lives," Cassandra Unchained, October 1, 2026 (free section); Stocktwits and Benzinga coverage, October 1–2, 2026.*
- *NVIDIA Form 10-Q for the quarter ended July 26, 2026 (commitments and guarantees notes, risk factors, Exhibit 10.1), signed August 26, 2026; NVIDIA Non-Deal Roadshow presentation, September 2026 (pp. 9, 10, 11, 21, 22); NVIDIA Q2 FY2027 earnings call, August 26, 2026; NVIDIA at Goldman Sachs Communacopia + Technology, September 10, 2026.*
- *Silicon Data GPU Forward Residual Value Methodology Overview, June 29, 2026; Silicon Data H100, A100 and B200 rental indices and residual values, accessed October 2, 2026; Silicon Data on X, August 11, 2026.*
- *Micron Q4 FY2026 earnings call, September 30, 2026; Samsung Electronics Q2 2026 earnings call, July 30, 2026; SK hynix Q2 2026 earnings call, July 29, 2026.*
- *CoreWeave Form S-1, March 3, 2025; CoreWeave Q2 2026 earnings call, August 11, 2026; CoreWeave at Goldman Sachs Communacopia, September 8, 2026; CoreWeave Form 8-K Exhibit 99.1 (DDTL 5.5 press release), August 10, 2026; JPMorgan, "CoreWeave: Upgrade to OW," September 24, 2026.*
- *Nebius Q2 2026 earnings call, August 12, 2026; Nebius press release, "Nebius raises $775 million in first secured debt financing to accelerate global buildout," July 17, 2026; Nebius at Goldman Sachs Communacopia, September 8, 2026.*

- *Yahoo Finance/Insider Monkey, "CoreWeave Pays SOFR +5.5%. Nebius Pays +2.5%," September 14, 2026; Quartz, July 31, 2026; CNBC, August 21, 2026; Financial Times, "Nvidia turns to insurers to spread risk of AI build-out," September 29, 2026; NVIDIA press release, "NVIDIA Partners With Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR…," August 10, 2026; Jensen Huang, "NVIDIA AI Factory Compute Is Becoming an Investable Asset Class," NVIDIA blog, August 11, 2026.*
- *Federal Reserve Bank of New York SOFR data, October 1, 2026; Markit five-year CDS for CoreWeave via LSEG (daily, September 30, 2025 to October 2, 2026); LSEG Workspace multiples and NVIDIA EBITDA estimates, October 2–3, 2026; Yahoo Finance historical prices.*
- *Adam Smith (George J. W. Goodman), The Money Game, 1968.*