When to Buy Is a Duration Question

I asked twenty-five value investors how they time their entries. The best answers were all describing the same variable, and none of them named it.

On September 22 I asked value investors on Substack a question in a Note: how do you use technical analysis. Twenty-five answered, and almost all gave a version of the same thing. Fundamentals set the value, charts time the entry.

That is right, and it stops one step short. The step it misses has a name, and the name is duration. I agreed with the consensus when I read it, then went back through my own five positions to see how I actually time entries. How much the entry even mattered turned out to depend almost entirely on one thing, and on most value names it matters far less than the people timing them think.

The short version

- The near-universal answer among value investors: fundamentals set the value, charts time the entry. It is right, and it stops one step short.
- Timing an entry only ever captures one part of a return. Your yield and your growth accrue with time no matter what day you buy. The only thing your entry price changes is the multiple you pay against the multiple you sell.
- How much that part is worth depends on how much of your return is cash paid to you no matter what, versus a future multiple the market may or may not grant. The more it leans on the multiple, the more timing matters.
- On my own five names the spread is not subtle. Buy on a random day and hold a year, and your outcome varied far more on one than the other: a standard deviation, the usual measure of how spread out the results are, of 18 points on Cheniere against 91 on NVIDIA.
- What would change my mind is below, and it is a number.

A return has three parts

Break a holding-period return into its pieces. You earn the starting cash yield. You earn the growth in cash flow per share over the years you hold. And you earn, or lose, the change in the multiple between the price you paid and the price you sold.

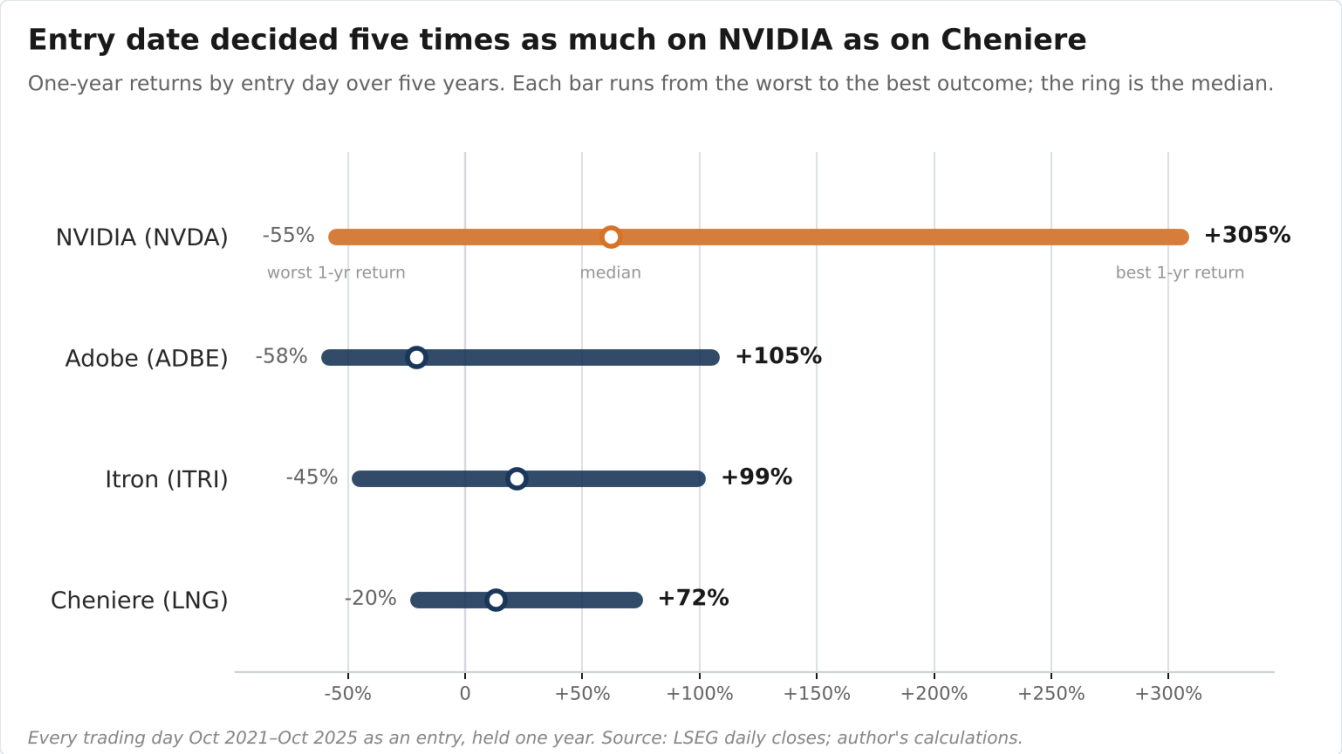
The first two accrue with time. Buy a compounder on a Tuesday or three weeks later and the yield and the growth barely notice. The third is the only piece your entry price controls. So timing an entry is a bet on one thing: the change in the multiple.

That reframes the whole question. Timing your entry is worth the effort in proportion to two things: how much of your expected return comes from the multiple changing, and how much that multiple swings around while you hold. A large, volatile multiple component means a good or bad entry moves your return a lot. A small, stable one means it barely registers.

How much the entry actually mattered

Before the theory, the measurement, because the measurement is what sent me looking for the theory. Take every trading day from October 2021 to October 2025, buy that day, hold one year on price alone, record the result. That is four years of entry dates, roughly a thousand of them, each held through a close in the last twelve months, drawn from five years of LSEG daily closes. The spread of those outcomes is how much the date you picked changed where you ended up.

Across those four years of entries the standard deviation of one-year price returns was 18 percentage points for Cheniere and 91 for NVIDIA. That is the measure, and only that: the historical dispersion of entry-cohort outcomes. It does not prove duration caused the gap, since NVIDIA is also simply a far more volatile stock, and I come back to that. But the ranking is the thing to sit with. The name whose return leans most on a future multiple is the one where the entry mattered most, and the name paid in contracted cash is the one where it was closest to noise.



One-year price return by entry date, four years of entries held one year each. Each bar runs from the worst outcome to the best; the wider the spread, the more your entry date decided the result.

The rest of this post is why that spread tracks what it tracks. The answer starts with duration.

What duration is

Duration means how far out a company's cash flows sit. The bond market has the cleanest version of it. A three-year zero and a thirty-year coupon bond are both bonds, but move the interest rate a point and the thirty-year swings far more, because more of what you are owed arrives further away. Stocks work the same way. A cheap, cash-generative name is the short bond: most of the value comes back to you soon, so a change in the discount rate barely moves it. A high-multiple grower is the long bond: most of what you are paying for arrives years out, past the end of any ten-year model, in the terminal value that the exit multiple stands in for. So its price moves more when the rate moves, and more of what you own is a number you have been promised, not yet paid.

One wrinkle, because Cheniere breaks the simple version. Its fee contracts run fifteen years, which by the textbook makes it long duration. It does not behave that way, because the cash is contracted and handed back along the way and a floor caps the repricing. So the axis that actually decides entry timing is not time-to-cash-flow. It is how much of your return is cash paid to you no matter what, against a future multiple the

market may or may not grant. Committed cash on one end, a discretionary multiple on the other. The high-multiple names sit at the discretionary end, and that is where the entry decides the most.

Duration is the biggest reason that share differs across businesses, but not the only one; volatility, leverage and sentiment can load the multiple too, which is why a name can swing more than its cash-flow profile alone would say. So two things set how much the entry matters, not one. The first is how much of your return rides on that discretionary multiple. The second is how little floor sits under a bad entry. A name can lean hard on the multiple and still be forgiving to enter if the downside is floored, which is exactly where Itron lands. The most entry-sensitive name is the one with both at once: its return rides on the multiple, and it has nothing underneath it.

I expected the textbook arithmetic to carry the argument, so I ran it. Take two companies worth the same today, both discounted at 9%, both growing for ten explicit years and then fading to 3% forever. One grows 3% the whole way, the other 15% through the decade with most of its value landing years out. Raise the discount rate a single point. The slow grower loses about 14% of its value, the fast grower about 17%. Wider, the way duration says it should be, but barely. That is the whole textbook effect, and everyone knows it. It is also not the whole story, because a rate scare drains sentiment out of a long-duration multiple on top of the math, and I think that second channel is the larger one. I cannot prove that from a DCF. I can show it on the tape, where my own names move far more with rates than 14 against 17 would predict.

One line of regime, because it raises the stakes. The US ten-year yield was 5.29% at the start of September and 4.96% by October, with the thirty-year still above 5% and the term premium rising, the backdrop from the gold post. A higher and jumpier discount rate makes the multiple on long-duration equities swing harder. That cuts both ways, and the second way is the opportunity. If the business is fine and only the rate moved, the swing hands you a better entry. Adobe fell to about $190 in June, on a management change and sour sentiment rather than the numbers, which looked fine. Buy it there and you are up roughly a third today, with the real question, the moat, no closer to an answer than it was. The gain came from the price swinging back, not from the business proving anything. The entry paid off before the thesis did.

My positions, in order of how much timing matters

Cheniere (LNG). The rare name that offers yield, contracted growth and a re-rating all at once, with real downside protection underneath. Roughly 90% of production is contracted through the mid-2030s on a weighted-average contract life near 15 years, part of $107.7B of fixed-fee revenue contracted in total, with about $6.6B of that scheduled for calendar 2026 regardless of commodity prices (FY2025 10-K). Those fees carry a high margin, because the variable part of each contract passes the feedgas cost through to the

customer, so most of the fixed fee drops toward cash earnings once fixed operating costs are covered. Take a conservative cut of it, call it $5.5 to $6 billion of contracted EBITDA, put it on 11x, strip $22.5 billion of net debt and $5.3 billion of CQP minority interest that Cheniere consolidates but does not fully own, across 206.5 million shares, and you land somewhere between about $160 and $185 against a price of $299. That is an illustrative downside on stated assumptions, not a hard floor the market owes you. (It is well below the ~$217 an earlier draft of this post showed, which capitalized the fee *revenue* directly at 11x and so quietly assumed a 100% margin; correcting that is most of the drop.) The protection I actually lean on is simpler and observed, not modelled. Over five years Cheniere's worst peak-to-trough fall was 26%, against 48% to 72% for the other three, and it returns close to 5% a year in cash, a 0.8% dividend plus a buyback that ran $1.1 billion in the first half, carrying the count from 207 million toward a 175 million target, none of which needs the multiple to cooperate. And it is cheap, about 10.6x next year's EBITDA (guided to $7.9 to $8.4 billion for 2026), because the market treats the war-driven margin spike as temporary. I think that underrates the compounding underneath: new trains and fresh long-term SPAs (QatarEnergy just signed for 4 mtpa through 2054, Petrobras for 0.8 mtpa) adding capacity, the 20/20 vision on DCF per share, the shrinking count. So this is committed cash and a low multiple at once. Its five-year beta to the market is 0.01. Timing matters least here, not for lack of upside but because the contracted cash and the capital returns cushion a bad entry more than anything else in the book.

Adobe (ADBE). Not a cash-flow risk. The revenue is recurring and the enterprise side is genuinely sticky, with switching costs around AEP and workflow lock-in. The duration here is in the narrative: what you are waiting on is not uncertain far-off cash flows, the cash is fine, but how long until the market's story about the moat resolves. The stock trades at 7.7x EV/EBITDA, below every large-cap software peer in the set and cheaper even than Oracle at 9.2x, against a peer median near 15.7x (Oracle, Intuit, Salesforce, Autodesk, Microsoft, SAP, ServiceNow). The question the market is actually asking is blunt: does AI kill Adobe? My read is that it is two businesses wearing one ticker. AI amplifies the enterprise half, because more AI-generated content means more need for Adobe's platform to manage, govern and deploy it, the same way more agents mean more need for a ServiceNow. And AI genuinely threatens the individual creative half, where open-source and AI-native tools are real substitutes. The market is pricing the blended company as if the whole thing is the creative half. The re-rate comes when that stops. The new CEO gets his first shot at reframing it on the December 9 call, his first as CEO, with the Q4 numbers underneath, net new ARR and the enterprise ARR trend most of all. Until then the cheapness is a bet on a story, and you cannot time a story.

NVIDIA (NVDA). The cheapest of the long-duration names, about 14x forward earnings and a low PEG alongside the rest of the chips, priced as if earnings are near a peak and the build-out is about to slow. At 14x the market is not paying for years of future demand, it is pricing a slowdown, so the upside is simply that the demand holds and the earnings keep growing into a multiple that is already compressing. The risk runs the

other way, and it is not demand rolling over first, the market is forward-looking enough to price that early. It is the memory shortage easing into oversupply, chip rents falling, and the neocloud financing chain that has been paying for all this cracking. The Burry cascade I wrote about. The reason the tape below shows the widest spread of outcomes by entry date of anything I own is its sheer volatility, not an expensive multiple.

Itron (ITRI). The business is compounding on its own, ARR up 21% and gross margin already past the 2027 target three years early. What sets it apart is where the two sides of the trade sit. The upside is almost all one thing: the market changing how it values the company at all, from an earnings multiple on what it still files under hardware, about 12.5x forward earnings today, toward what it would pay for the software business Itron is becoming. The downside is floored from the other side, because I am buying cheap even on the current, conservative framework, and the demand underneath is not discretionary. Tom Deitrich, the CEO, made the point at a Roth fireside chat: the grid build-out is inevitable, the only real rate limiter is how fast utilities adopt the technology. "It is going to happen," he said. When, not if. So the software mix-shift accrues underneath regardless, and being early costs me time, not capital. The re-rate itself waits on the Investor Day in early 2027, when the model rebuild gets hard to avoid.

Gold. Not the far end of this list. Off it entirely. Gold has no corporate cash flows and no conventional valuation multiple, so the framework I use for equities largely breaks down. The textbook would call a non-yielding asset the longest duration there is, and before 2022 that held: Western funds traded gold against real yields, and it fell when rates rose. That link broke. The marginal buyer is now central banks, who bought a record pace from 2022 through 2025 while yields climbed, and gold rallied into rising rates rather than falling with them. So I do not try to time it. I accumulate against a floor the central-bank buyer keeps lifting, and I watch the break conditions. I hold it mostly through Wheaton, which does have cash flows and a multiple and so does sit on the scale above, with a smaller GLD position as pure-metal exposure, to hedge the company-specific risk a streamer carries on top of the gold price.

The Itron day that proves it

On July 28 Itron rose 26% in a single session on a margin beat and an EPS guide. I wrote at the time that the move repriced the EPS and did not reprice the business. In this framework that sentence has a precise meaning.

The 26% was the market capitalizing a higher expected-earnings base, not awarding a re-rate. Investors marked up what they thought Itron would earn, then kept paying the same kind of multiple for it, an earnings multiple on a business they still file under hardware. You can see it in the multiple itself: even after the 26% jump the stock sat at about 15.6x forward earnings, 0.75x the S&P against a 1.29x ten-year average, near the

trough of its own range. The market marked the earnings up and kept paying a trough, hardware-style multiple for them. A name that used to trade at a 29% premium to the market was still at a 25% discount.

Here is what happened next, and it is the whole argument in one name. The stock has given the entire move back. Itron closed at $87.65 on October 23, below the $99.66 it settled at after the print and below my own cost, and cheaper on forward earnings than it was before the quarter that beat. It now sits about 4% below its 200-week moving average. The earnings pop faded because the multiple never re-rated, and the re-rate, the switch to valuing the software, is still sitting out there waiting on a catalyst with a date. Being early has cost me a few percent, and nothing of the thesis: the piece I am waiting on has not moved.



Itron gave the entire 26% back and now sits on its 200-week line, with the re-rate still ahead.

Where I actually sit

My own entries against the October 23 close, for honesty about the cost of my own timing:

Name	my entry	now	change
Cheniere (LNG)	272.03	298.67	+9.8%
Adobe (ADBE)	274.14	263.82	−3.8%
NVIDIA (NVDA)	179.78	210.96	+17.3%
Itron (ITRI)	92.74	87.65	−5.5%

The two I am down on, Adobe and Itron, are the two whose return leans hardest on a re-rate that has not happened yet. The two I am up on are carried by things that accrue no matter when I bought: Cheniere by its contracted cash and the cushion under it, NVIDIA by the growth that has already shown up. That is the framework describing my own book back to me, before I run a single regression on it.

I am acting on it as I write, and the framework is changing my sizing, not keeping score. This week I trimmed Cheniere and took a little off Adobe ahead of a December binary I cannot handicap, and moved both into gold. Not because the framework called a top, but because it told me how much less of my return on those two was ever going to come from an entry I could time, against cash I am paid regardless and downside that does not need the market to agree.

What actually makes a catalyst

A credible, dated catalyst shortens the wait for a re-rating. It does not change the duration of the underlying cash flows; it pulls the re-rate forward in time, which raises what that piece is worth per year and cuts how long your capital sits idle. But a date on the calendar is not a catalyst on its own, and this is the part most timing talk skips.

In a name everyone watches, you will not out-read the next earnings print. The bar is too well set, and if the market already expects a good print, a good print changes nothing. The expectation is in the price, the stock needs a beat, and beating a bar the market set high is hard. So the edge is rarely in the earnings print itself for a watched name. It is in a view the consensus does not hold yet, plus a reason it cannot hold it. The events that actually re-rate a stock do two things at once. They force the consensus to re-classify the business, and they remove the specific reason the consensus had not done it already. That second part is the one most timing talk skips, because if there were no reason the market was wrong, it would not be wrong. With Itron the reason is concrete: most of the analysts covering it are industrials people without a software background, and none of them can build an EV/ARR model on a segment the company only began breaking out three quarters ago, with no long-range framework to anchor it. The Investor Day is the catalyst precisely because it

hands them that framework. In the meantime the software segment growing faster than the rest, quarter after quarter, is the running signal that the reclassification is coming. So the test is not whether something good happens, which is usually priced, but whether it removes the barrier that kept your reading from becoming the consensus one. Itron reporting a strong quarter is priced. Itron giving analysts the segment framework to rebuild the model on EV/ARR is not, because the market is narrative-heavy and will not adopt your reading of a business until the evidence leaves it no choice. That is the second-level version of a catalyst. Not what happens, but what the market is made to admit it had wrong, once it can no longer explain the gap away.

Sizing falls out of it, roughly and all else equal. A wide gap to fair value plus a short, dated catalyst earns a larger position, though a near-dated catalyst also concentrates event risk, so downside and portfolio weight still set the size; the framework tells you how much the entry matters, not the whole position. A wide gap with no clock, Adobe or gold, earns patience, not size. How you tell a real catalyst from a priced-in one before the move, and buy ahead of it, is the harder question, and its own post.

FOMO is a duration error

FOMO is a duration error before it is a feeling. The mistake is specific: you pay a peak multiple for the longest-dated, most multiple-dependent cash flows at the worst possible moment, right after the re-rate has already happened for someone else. Howard Marks put the other side of it plainly: "Everyone hopes a buyer will come along who's willing to overpay for what they have for sale. But certainly the hoped-for arrival of this sucker can't be counted on." Buying a long-duration name after its multiple has run is a bet that the sucker shows up.

It hurts most on the long-duration names, because that is where the multiple falls furthest when the rate normalises or the story cracks. The October 8 selloff showed it cleanly. NVIDIA fell about 10% in the days after. CoreWeave and Nebius, the leveraged neoclouds whose whole value is a bet on future compute demand, fell about 31% and 30%, three times as far, on the same news. Leverage did some of that work, since the neoclouds carry debt NVIDIA does not, so this is the phenomenon on display rather than clean proof. But the ranking is the one the rest of this post keeps finding: the longest-duration, most multiple-dependent names took the deepest cut.

The rest of the tape

The spread at the top of this post was the first angle. Here are two more from the same five years of closing prices, and neither needs a regression to read.

Drawdowns say the same thing, and they need no model at all. Not the one-year windows above, just the worst peak-to-trough fall each name actually suffered over the five years: Cheniere −26%, Itron −48%, NVIDIA −66%, Adobe −72%. The floored name fell least. The names whose return rides on a multiple fell two to three times as far. That Cheniere, the one carried by contracted cash, took the shallowest fall is the whole thesis in a single number, and it needs no regression to see.

Rate sensitivity points the same way. Regress each stock's monthly return on the monthly move in the ten-year yield over the last five years and the ranking holds. Cheniere barely reacts, and the long-duration names move about ten times as hard.

Sensitivity to the 10-Year Yield

Name	Move per +1 point in the 10-year	R^2
Cheniere (LNG)	−1.6%	0.00
Itron (ITRI)	−10.6%	0.08
Adobe (ADBE)	−15.7%	0.16
NVIDIA (NVDA)	−16.2%	0.10

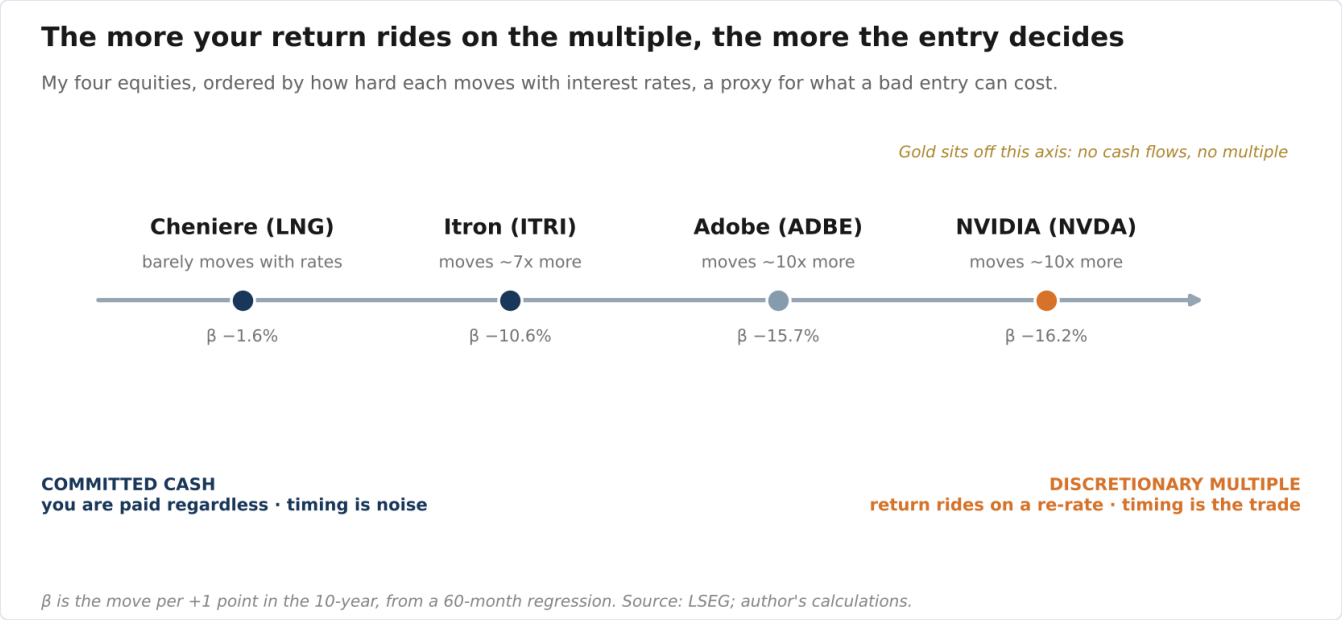
60-month regression of monthly returns on the monthly change in the 10-year yield, Oct 2021–Oct 2026. Source: LSEG; author's calculations.

How hard each name moves with the 10-year. The floored name barely reacts; the re-rate names move about ten times as much.

The R-squared runs from roughly zero on Cheniere, where rates explain almost none of its moves, to 0.16 on Adobe, so read the ranking, not the decimals. The drawdown above says the same thing without any regression at all, which is why I lead with it.

This is the evidence for the claim I could only assert earlier. The textbook DCF said a one-point rate rise should cost a long-duration name about three points more than a short one, 17 against 14. The tape says the long-duration names move something like ten times as much with rates as Cheniere does. The cash-flow forecast is the same in both accounts, so the gap between what the model says and what the price does is, by construction, the multiple moving. What I cannot do is pin that move cleanly on investors repricing the far-out years. Rates, risk appetite, factor positioning and plain sentiment all run through these high-beta names at

once, and the regression cannot separate them. So I will claim only what the ranking supports: the more multiple-dependent names carry more rate sensitivity, and the multiple is doing most of the work the arithmetic never captured.



My four equities, ordered by what a bad entry can cost. The entry matters more as you move right.

This chart orders by the second of the two legs from earlier, the floor. The rate beta measures how hard the multiple gets marked down when money turns, which is what decides whether a bad entry actually costs you, so NVIDIA, a volatile multiple with no floor, sits at the far end and Cheniere at the near one. Itron sits in the middle for a real reason, not a measurement quirk: its downside is floored, because I am cheap on the current framework and the demand underneath is not discretionary. What the beta cannot show is Itron's first leg, the share of its upside riding on a re-rate that is still ahead, and company-specific rather than rate-driven. That is the Itron section above. The spread of one-year outcomes at the top of the post, which does not care why a multiple moved, is the cleaner test, and it is the one I lean on.

One last cut, because it is the rule the chart crowd actually uses. There is a well-worn line, usually hung on Charlie Munger though the quote is almost certainly apocryphal, that you buy quality names when they fall to their 200-week moving average. The 200-week part matters, and it is the part people get wrong by reaching for the 200-day, as I did at first. A quality name touches its 200-day line constantly. It touches its 200-week line only in a real dislocation. Run the 200-week line on the four names with price histories and it flags exactly one. Itron sits about 4% below its 200-week line today. Cheniere trades 44% above its line and NVIDIA 77% above, and neither has been near it in five years. Adobe sits 34% below a 200-week line dragged up by three years of $400-to-$600 prices that are not coming back, a lagging average, not a buy signal. The rule points at the one name where the re-rate has not happened, and says nothing on the names

carried by growth or capped by a floor. One firing across five years is an observation, not a test. A real test of whether the rule works wants fifteen years and a cross-section of names, and that is a separate post.

Read against this, the survey answers stop arguing. The investors who use charts for entries are reaching for the same thing: a read on when the multiple, the sentiment piece, has swung too low and is likely to snap back. Those are the mean-reverting re-rate names, where the multiple oscillates around a level and buying it depressed catches the bounce, the Itrons and the Adobes. It adds much less on a short-duration cash machine, a tobacco name, a regulated utility, a contracted energy name like Cheniere, where the multiple is a small, stable slice of the return, so it swings less and the entry matters less. Not nothing, just less. The investor who ignores charts entirely is close to right for that book. The two camps are not disagreeing about technical analysis. They are describing different durations.

What would change my mind

The framework makes one prediction: entry timing should matter more on the names where the return rides on a multiple than on the ones paid in cash. The spread of one-year outcomes at the top is where that gets tested. Had Cheniere's one-year outcomes varied as much by entry date as NVIDIA's, the prediction would be false and the framework with it. Instead Cheniere's spread is a fifth of NVIDIA's. And a chart alone is never a reason to buy: it only tells you whether the multiple is depressed, which matters once you have already decided the business is worth owning.

The close

The chart was never telling me when to buy. It was telling me how much of the return rode on the multiple. That was a duration question the whole time.

So the question is not whether to time your entries. It is which of your positions have earned the effort. Pull up your own book and sort it: the holdings where the return is cash you will be paid regardless, and the ones where it rides on the market changing its mind. Then ask how many you have been timing out of habit, when the return was never going to come from the multiple in the first place. Most people run it backwards.

Disclosure: I hold positions in Cheniere (LNG), Adobe (ADBE), NVIDIA (NVDA), Itron (ITRI), Wheaton Precious Metals (WPM) and the SPDR Gold Shares ETF (GLD). This material is for informational purposes only and does not constitute investment advice. Sources:

- *Price and yield data: LSEG daily closing prices and US 10-year yield, October 2021 to October 2026; dispersion, drawdown and rate-beta calculations are the author's own. Returns are price returns, with dividends excluded as immaterial at these yields.*

The entry-cohort dispersion uses entries from October 2021 to October 2025, each held one year. The Cheniere downside figure is an illustrative valuation scenario on stated assumptions, not a hard floor.

- *The September 22, 2026 Note on technical analysis among value investors on Substack.*
- *Cheniere: FY2025 Form 10-K (contracted SPA revenue table); Q2 2026 earnings call and press release, August 6, 2026; QatarEnergy SPA press release, October 22, 2026; Petrobras SPA press release, September 29, 2026; current multiples, net debt, share count and dividend via TIKR Terminal and StockAnalysis.*
- *Itron: Q2 2026 earnings call, July 28, 2026; Roth MKM fireside chat with CEO Tom Deitrich and CFO Joan Hooper, August 5, 2026; and prior coverage ("My Biggest Position Went Up 26% in One Day," August 3, 2026).*
- *Adobe: Q3 FY2026 earnings call, September 10, 2026; forward EV/EBITDA multiples for ADBE and peers (CRM, ADSK, INTU, SAP, MSFT, ORCL, NOW) via LSEG Workspace; prior coverage ("Is Adobe a Value Trap?," September 14, 2026).*
- *NVIDIA and the neoclouds: prior coverage ("Burry Is Right About the GPUs. He's Wrong About NVIDIA.," October 5, 2026); October 8 selloff prices via Yahoo Finance.*
- *Gold: World Gold Council central-bank demand data; prior coverage ("The Current Thesis on Gold," September 8, 2026).*
- *Howard Marks, Oaktree memo. Equity-duration arithmetic is the author's own ten-year discounted cash flow.*
